

LU-17: INFERENCE IN FIRST ORDER LOGIC

LU Objectives

To explain differed inference mechanisms

To study first order resolution technique

LU Outcomes

CO : 3

Apply inference rules

Implement automated theorem provers using resolution mechanisms

Outline

- Reducing first-order inference to propositional inference
- Unification
- Generalized Modus Ponens

Universal instantiation (UI)

- Every instantiation of a universally quantified sentence is entailed by it:
-

$$\frac{\forall v \alpha}{\text{Subst}(\{v/g\}, \alpha)}$$

for any variable v and ground term g

- E.g., $\forall x \text{ King}(x) \wedge \text{Greedy}(x) \Rightarrow \text{Evil}(x)$ yields:
-
- To write out the inference rule formally, we use the notion of **substitutions**
 $\text{King}(\text{John}) \wedge \text{Greedy}(\text{John}) \Rightarrow \text{Evil}(\text{John})$
 $\text{King}(\text{Richard}) \wedge \text{Greedy}(\text{Richard}) \Rightarrow \text{Evil}(\text{Richard})$
 $\text{King}(\text{Father}(\text{John})) \wedge \text{Greedy}(\text{Father}(\text{John})) \Rightarrow \text{Evil}(\text{Father}(\text{John}))$
- For example, the three sentences given earlier are obtained with the substitutions $\{x/\text{John}\}$, $\{x/\text{Richard}\}$, and $\{x/\text{Father}(\text{John})\}$.

Existential instantiation (EI)

- For any sentence α , variable v the variable is replaced by constant symbol k that does **not** appear elsewhere in the knowledge base:

-

$$\frac{\exists v \alpha}{\text{Subst}(\{v/k\}, \alpha)}$$

- E.g., $\exists x \text{Crown}(x) \wedge \text{OnHead}(x, \text{John})$ we can infer

$$\text{Crown}(C_1) \wedge \text{OnHead}(C_1, \text{John})$$

provided C_1 is a new constant symbol, called a **Skolem constant**

Reduction to propositional inference

Suppose the KB contains just the following:

$\forall x \text{ King}(x) \wedge \text{Greedy}(x) \Rightarrow \text{Evil}(x)$
King(John)
Greedy(John)
Brother(Richard, John)

- Instantiating the universal sentence in **all possible** ways, we have:
King(John) \wedge Greedy(John) \Rightarrow Evil(John)
King(Richard) \wedge Greedy(Richard) \Rightarrow Evil(Richard)
King(John)
Greedy(John)
Brother(Richard, John)
- The new KB is **propositionalized**: proposition symbols are
-

King(John), Greedy(John), Evil(John), King(Richard), etc.

Reduction contd.

- Every FOL KB can be propositionalized so as to preserve entailment
-
- (A ground sentence is entailed by new KB iff entailed by original KB)
-
- Idea: propositionalize KB and query, apply resolution, return result
-
- Problem: with function symbols, there are infinitely many ground terms,
 - e.g., *Father(Father(Father(John)))*
 -

A first-order inference rule

- E.g:
 $\forall x \text{ King}(x) \wedge \text{Greedy}(x) \Rightarrow \text{Evil}(x)$
 $\text{King}(\text{John})$
 $\forall y \text{ Greedy}(y)$
 $\text{Brother}(\text{Richard}, \text{John})$
- We would still like to be able to conclude that $\text{Evil}(\text{John})$, because we know that John is a king (given) and John is greedy (because everyone is greedy).
- We have to find a substitution both for the variables in the implication sentence and for the variables in the sentences that are in the knowledge base.
- By applying the substitution $\{x/\text{John}, y/\text{John}\}$ to the implication premises $\text{King}(x)$ and $\text{Greedy}(x)$ and the knowledge-base sentences $\text{King}(\text{John})$ and $\text{Greedy}(y)$ will make them identical. Thus, we can infer the conclusion of the implication.
- This inference process can be captured as a single inference rule is called **Generalized Modus Ponens**

A first-order inference rule

- For atomic sentences p_i' , p_i , and q , where there is a substitution θ such that $\text{SUBST}(\theta, p_i') = \text{SUBST}(\theta, p_i)$, for all i ,

$$\frac{p_1', p_2', \dots, p_n', (p_1 \wedge p_2 \wedge \dots \wedge p_n \Rightarrow q)}{\text{SUBST}(\theta, q)}.$$

- There are $n+1$ premises to this rule: the n atomic sentences p_i' and the one implication.
- The conclusion is the result of applying the substitution θ to the consequent q .

A first-order inference rule

- Ex:

$p1'$ is King(John)

$p1$ is King(x)

$p2'$ is Greedy(y)

$p2$ is Greedy(x)

θ is $\{x/\text{John}, y/\text{John}\}$

q is Evil(x)

$\text{SUBST}(\theta, q)$ is Evil(John)

- This Generalized Modus Ponens is a sound inference rule.
- For any sentence p (whose variables are assumed to be universally quantified) and for any substitution θ ,

$$p \models \text{SUBST}(\theta, p)$$

- Thus, from $p1', \dots, pn'$ we can infer

$$\text{SUBST}(\theta, p1') \wedge \dots \wedge \text{SUBST}(\theta, pn')$$

and from the implication $p1 \wedge \dots \wedge pn \Rightarrow q$ we can infer

$$\text{SUBST}(\theta, p1) \wedge \dots \wedge \text{SUBST}(\theta, pn) \Rightarrow \text{SUBST}(\theta, q)$$

A first-order inference rule

- θ in Generalized Modus Ponens is defined so that $\text{SUBST}(\theta, p_i') = \text{SUBST}(\theta, p_i)$, for all i ;
- Therefore the first of these two sentences matches the premise of the second exactly. Hence, $\text{SUBST}(\theta, q)$ follows by Modus Ponens.
- Generalized Modus Ponens is a **lifted version of Modus Ponens**. It raises Modus Ponens from ground (variable-free) propositional logic to first-order logic.

Problems with propositionalization

- Propositionalization seems to generate lots of irrelevant sentences.
- E.g., from:
 -
 -
 - $$\forall x \text{ King}(x) \wedge \text{Greedy}(x) \Rightarrow \text{Evil}(x)$$
$$\text{King}(\text{John})$$
$$\forall y \text{ Greedy}(y)$$
$$\text{Brother}(\text{Richard}, \text{John})$$
- it seems obvious that *Evil(John)*, but propositionalization produces lots of facts such as *Greedy(Richard)* that are irrelevant
 -
- With p k -ary predicates and n constants, there are $p \cdot n^k$ instantiations.
 -

Unification

- Lifted inference rules require finding substitutions that make different logical expressions look identical. This process is called **unification**
- Unification and is a key component of all first-order inference algorithms.
- The UNIFY algorithm takes two sentences and returns a **unifier** for them if one exists:

$$\text{UNIFY}(p, q) = \theta \text{ where } \text{SUBST}(\theta, p) = \text{SUBST}(\theta, q) .$$

Unification

- Ex: We have a query
AskVars(Knows(John, x))
- This query means “whom does John know?”. Answers to this query can be obtained by finding all sentences in the knowledge base that unify with Knows(John, x).
- The results of unification with four different sentences that are in the knowledge base:
 - $\text{UNIFY}(\text{Knows}(\text{John}, x), \text{Knows}(\text{John}, \text{Jane})) = \{x/\text{Jane}\}$
 - $\text{UNIFY}(\text{Knows}(\text{John}, x), \text{Knows}(y, \text{Bill})) = \{x/\text{Bill}, y/\text{John}\}$
 - $\text{UNIFY}(\text{Knows}(\text{John}, x), \text{Knows}(y, \text{Mother}(y))) = \{y/\text{John}, x/\text{Mother}(\text{John})\}$
 - $\text{UNIFY}(\text{Knows}(\text{John}, x), \text{Knows}(x, \text{Elizabeth})) = \text{fail}$.
- The last unification fails because x cannot take on the values John and Elizabeth at the same time.

Unification

- WKT that $\text{Knows}(x, \text{Elizabeth})$ means “Everyone knows Elizabeth,”
- So we *should be able to infer that John knows Elizabeth*.
- The problem arises only because the two sentences happen to use the same variable name, x .
- The problem can be avoided by **standardizing apart** one of the two sentences being unified, which means renaming its variables to avoid name clashes.
- For example, we can rename x in $\text{Knows}(x, \text{Elizabeth})$ to $x17$ (a new variable name) without changing its meaning. Now the unification becomes

$$\text{UNIFY}(\text{Knows}(\text{John}, x), \text{Knows}(x17, \text{Elizabeth})) = \{x/\text{Elizabeth}, x17/\text{John}\} .$$

Unification

- There is one more complication.
- Generally UNIFY should return a substitution that makes the two arguments look the same. But there could be more than one such unifier.
- For example, $\text{UNIFY}(\text{Knows}(\text{John}, x), \text{Knows}(y, z))$ could return $\{y/\text{John}, x/z\}$ or $\{y/\text{John}, x/\text{John}, z/\text{John}\}$.
- The first unifier gives $\text{Knows}(\text{John}, z)$ as the result of unification, whereas the second gives $\text{Knows}(\text{John}, \text{John})$.
- The first unifier is *more general than the second, because it places fewer restrictions on the values of the variables*.
- For every unifiable pair of expressions, there is a single **Most General Unifier (MGU)** that is unique up to renaming and substitution of variables.

The unification algorithm

```
function UNIFY( $x, y, \theta$ ) returns a substitution to make  $x$  and  $y$  identical
  inputs:  $x$ , a variable, constant, list, or compound
            $y$ , a variable, constant, list, or compound
            $\theta$ , the substitution built up so far

  if  $\theta = \text{failure}$  then return failure
  else if  $x = y$  then return  $\theta$ 
  else if VARIABLE?( $x$ ) then return UNIFY-VAR( $x, y, \theta$ )
  else if VARIABLE?( $y$ ) then return UNIFY-VAR( $y, x, \theta$ )
  else if COMPOUND?( $x$ ) and COMPOUND?( $y$ ) then
    return UNIFY(ARGS[ $x$ ], ARGS[ $y$ ], UNIFY(OP[ $x$ ], OP[ $y$ ],  $\theta$ ))
  else if LIST?( $x$ ) and LIST?( $y$ ) then
    return UNIFY(REST[ $x$ ], REST[ $y$ ], UNIFY(FIRST[ $x$ ], FIRST[ $y$ ],  $\theta$ ))
  else return failure
```

The unification algorithm

```
function UNIFY-VAR( $var, x, \theta$ ) returns a substitution
  inputs:  $var$ , a variable
          $x$ , any expression
          $\theta$ , the substitution built up so far

  if  $\{var/val\} \in \theta$  then return UNIFY( $val, x, \theta$ )
  else if  $\{x/val\} \in \theta$  then return UNIFY( $var, val, \theta$ )
  else if OCCUR-CHECK?( $var, x$ ) then return failure
  else return add  $\{var/x\}$  to  $\theta$ 
```

Unification Algorithm

- The process of Unification is simple:
 - recursively explore the two expressions simultaneously “side by side,” building up a unifier along the way,
 - It fails if two corresponding points in the structures do not match.
- There is one expensive step: when matching a variable against a complex term, one must check whether the variable itself occurs inside the term;
- if it does, the match fails because no consistent unifier can be constructed.
- For example, $S(x)$ can't unify with $S(S(x))$.
- This is called **occur check**, it makes the complexity of the entire algorithm quadratic in the size of the expressions being unified

Storage and retrieval

- Underlying the TELL and ASK functions used to inform and interrogate a knowledge base are the more primitive STORE and FETCH functions
 - STORE(s) stores a sentence s into the knowledge base
 - FETCH(q) returns all unifiers such that the query q unifies with some sentence in the knowledge base.
- The simplest way to implement STORE and FETCH is to keep all the facts in one long list and unify each query against every element of the list. Such a process is inefficient, but it works.

Storage and retrieval


- FETCH can be more efficient by ensuring that unifications are attempted only with sentences that have *some chance of unifying*.
- *For example, there is no point in trying to unify Knows(John, x) with Brother (Richard, John). We can avoid such unifications by **indexing the facts in the knowledge base**.*
- **Predicate indexing** puts all the Knows facts in one bucket and all the Brother facts in another. The buckets can be stored in a hash table for efficient access.
- Predicate indexing is useful when there are many predicate symbols but only a few clauses for each symbol. Not suitable for a predicate has many clauses.

Storage and retrieval

- Ex: A predicate
 Employs(x, y)
- would have a very large bucket with perhaps millions of employers and tens of millions of employees.
- For this particular query, it would help if facts were indexed both by predicate and by second argument, perhaps using a combined hash table key
- For other queries, such as Employs(IBM, y), we would need to have indexed the facts by combining the predicate with the first argument.
- So facts can be stored under multiple index keys, so that they can be instantly accessible to various queries that they might unify with.


Storage and retrieval

- For the fact $\text{Employs}(\text{IBM}, \text{Richard})$, the queries are
 - $\text{Employs}(\text{IBM}, \text{Richard})$ Does IBM employ Richard?
 - $\text{Employs}(x, \text{Richard})$ Who employs Richard?
 - $\text{Employs}(\text{IBM}, y)$ Whom does IBM employ?
 - $\text{Employs}(x, y)$ Who employs whom?
- These queries form a **subsumption lattice**



Storage and retrieval

- For example, the child of any node in the lattice is obtained from its parent by a single substitution; and the “highest” common descendant of any two nodes is the result of applying their most general unifier.
- A sentence with repeated constants has a slightly different lattice,



Storage and retrieval

- Works very well whenever the lattice contains a small number of nodes.
- For a predicate with n arguments, however, the lattice contains $O(2^n)$ nodes.
- If function symbols are allowed, the number of nodes is also exponential in the size of the terms in the sentence to be stored.
- This can be solved by adopting a fixed policy, such as maintaining indices only on keys composed of a predicate plus each argument, or by using an adaptive policy that creates indices to meet the demands of the kinds of queries being asked.